

Locally run k6 load tested results in Grafana:



Terminal k6 load testing results:

TOTAL RESULTS

```
checks_total.....: 24579  204.254592/s
checks_succeeded...: 100.00% 24579 out of 24579
checks_failed.....: 0.00%  0 out of 24579
```

```
✓ POST status is 201
✓ GET status is 200
✓ DELETE status is 204
```

HTTP

```
http_req_duration.....: avg=1.97ms min=376µs    med=1.63ms max=304.98ms p(90)=3.23ms p(95)=3.85ms
  { expected_response:true }...: avg=1.97ms min=376µs    med=1.63ms max=304.98ms p(90)=3.23ms p(95)=3.85ms
http_req_failed.....: 0.00%  0 out of 24579
http_reqs.....: 24579  204.254592/s
```

EXECUTION

```
iteration_duration.....: avg=1.1s   min=711.19ms med=1.1s   max=1.64s   p(90)=1.32s p(95)=1.38s
iterations.....: 8193  68.084864/s
vus.....: 1    min=1    max=100
vus_max.....: 100  min=100  max=100
```

NETWORK

```
data_received.....: 3.8 MB 31 kB/s
data_sent.....: 3.0 MB 25 kB/s
```

Local Kind k8 grafrana.



Local SRE incident simulation:

Pod Crash — forces a running pod to die instantly, mimicking a server crash or process failure. Kubernetes notices it's gone and automatically starts a replacement. No human intervention needed.

